

1 **PhantomHunter: Detecting Unseen Privately-Tuned** 2 **LLM-Generated Text via Family-Aware Learning**

6 **Abstract**

7 With the popularity of large language models (LLMs), undesirable
8 societal problems like misinformation production and academic
9 misconduct have been more severe, making LLM-generated text
10 detection now of unprecedented importance. Though existing meth-
11 ods have made remarkable progress, they mostly consider publicly
12 known LLMs when testing the performance, and a new challenge
13 brought by text from privately-tuned LLMs is largely underexplored.
14 Due to the rapid development of open-source models like LLaMA
15 and Qwen series and efficient LLM training methods, even ordinary
16 users can now easily possess private LLMs by fine-tuning an open-
17 source one with private corpora. This could lead to a significant
18 performance drop of existing detectors in practice, due to their poor
19 capability of capturing the essential LLM traits robust to fine-tuning
20 operations. Our preliminary examination reveals that fine-tuning
21 an LLM with 11M tokens could make a detector's accuracy jump
22 from 100% to only 59% at most. To address this issue, we propose
23 PhantomHunter, an LLM-generated text detector specialized for
24 detecting text from unseen privately-tuned LLMs, whose family-
25 aware learning framework captures family-level traits shared across
26 the base models and their derivatives, instead of memorizing indi-
27 vidual characteristics. Specifically, PhantomHunter first extracts
28 base model features and enhances the family-shared information
29 using a contrastive family-aware learning module. The enhanced
30 features are then fed into a mixture-of-experts module containing
31 multiple experts for corresponding families for final predictions. Ex-
32 periments on data from four widely-adopted LLM families (LLaMA,
33 Gemma, Mistral, and Qwen) show its superiority over 8 baselines
34 and 4 industrial services.

36 **Keywords**

37 AI-generated Text Detection, Large Language Model, Family-aware
38 Learning

40 **ACM Reference Format:**

41 . 2026. *PhantomHunter: Detecting Unseen Privately-Tuned LLM-Generated*
42 *Text via Family-Aware Learning*. In *Proceedings of ACM Web Conference 2026*
43 (*WWW '26*). ACM, New York, NY, USA, 11 pages. <https://doi.org/XXXXXXXXXXXX>.

45 **1 Introduction**

47 Large language models (LLMs) have successfully revolutionized
48 the way people organize and produce text and inspire productivity

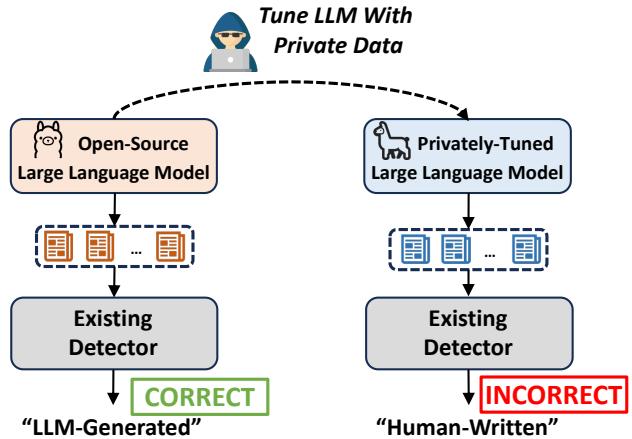
50 Permission to make digital or hard copies of all or part of this work for personal or
51 classroom use is granted without fee provided that copies are not made or distributed
52 for profit or commercial advantage and that copies bear this notice and the full citation
53 on the first page. Copyrights for components of this work owned by others than the
54 author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or
republish, to post on servers or to redistribute to lists, requires prior specific permission
and/or a fee. Request permissions from permissions@acm.org.

55 *WWW '26, Dubai, United Arab Emirates*

56 © 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

57 ACM ISBN 978-x-xxxx-xxxx-x/YY/MM

58 <https://doi.org/XXXXXXXX.XXXXXXXX>



80 **Figure 1: Illustration of how private fine-tuning of LLMs**
81 **affects LLM-generated text detection.** The existing detector
82 may successfully identify text generated by the original open-
83 source LLM, but fails to detect text generated by the privately
84 fine-tuned one due to its underlying reliance on the non-
85 shared features among the base LLM and the fine-tuned ones,
86 leading to missed detection.

88 in a wide range of applications [22]. However, undesirable societal
89 problems caused by the misuses of LLMs also rapidly emerge, such
90 as cheating in academic writing [23], creating misinformation [5],
91 and accelerating information pollution [34]. As the front-line barrier
92 against such threats, automatic detection techniques of LLM-
93 generated text (LLMGT) is now of unprecedented importance and
94 attracts researchers from both academia and industry.

95 LLMGT detection generally distinguishes LLM-generated and
96 human-written text via binary classification. Existing methods ei-
97 ther learn common textual features (e.g., stylistic clues [13]) shared
98 across LLMs using representation learning or design distinguishable
99 metrics between human and LLM texts based on LLMs' internal
100 signals (e.g., token probabilities [38]). For both categories, their
101 tests were mostly conducted on data from publicly available LLMs,
102 assuming that users generate text using public, off-the-shelf ser-
103 vices. **We argue that this situation is being changed due to**
104 **the recent development of the open-source LLM community.**
105 With the help of platforms like HuggingFace¹ and the efficient
106 LLM training techniques like low-rank adaptation (LoRA) [19],
107 building fine-tuned LLMs with customized private datasets has
108 become much easier than before. For instance, there have been
109 over 140k Qwen-based and 60k Llama-based derivative models on
110 HuggingFace [1, 30]. As shown in figure 1, after private fine-tuning
111 on unknown corpus, the learned characteristics of base models
112

113 ¹<https://huggingface.co/models>

could change and the LLMGT detectors would fail (will empirically show in the following section), shaping a new risk that malicious users can generate harmful texts privately without being caught by LLMGT detectors. A new challenge arises: **How could we detect text generated by privately-tuned open-source LLMs?**

To address this issue, we first conduct an analysis of differences between the base LLMs and its derivative models via fine-tuning on data of different topics. The result reveals that compared to other non-homologous models, the token probability lists of text on the fine-tuned LLM have a higher similarity to the probability lists on the base LLM, exhibiting clear “family trait”. Inspired by this finding, we propose to perform family-aware learning by extracting probabilistic features of base models and design **PhantomHunter**, a detector targeted at text generated by unseen privately-tuned LLMs. PhantomHunter consists of three main components: First, it obtains probability features from widely-known base LLMs such as Qwen and LLaMA. The extracted features are then enhanced using a contrastive family-aware learning to better preserve the family-level traits shared across the text pieces from both the base model and the fine-tuned derivative models. The enhanced features are fed into a Mixture-of-Experts (MoE) module that contains multiple expert module specialized in detecting texts from specific LLM families. Finally, the collective judgment of the given text begin LLM-generated or not is made by the gate mechanism controlled by the family prediction results. Experiments against baseline methods from both previous literature and the real-world industrial services demonstrates that our proposed PhantomHunter could better capture family-level traits and effectively detect generated text from unseen privately-tuned LLMs. Our main contributions are as follows:

- **New Risk:** We empirically expose the new risk that current LLM-generated text detectors may fail to detect text from privately-tuned LLMs.
- **New Method:** We propose PhantomHunter, which trained via family-aware learning to learn common traits shared in typical open-source LLM families.
- **High Performance:** We construct a new dataset containing texts in academic abstracts and Q&A genres generated by human and four widely-adopted LLM families (LLaMA, Gemma, Mistral, and Qwen) for evaluation. Experimental results show that PhantomHunter outperforms 7 baselines and 4 industrial services.

2 How Does Fine-tuning Affect LLM-Generated Texts Detection?

2.1 Preliminary Experiment

To explore the extent to which fine-tuning affects the performance of LLM-generated text detectors, we fine-tuned LLaMA-2 7B using arXiv abstract corpus by saving checkpoints at different fine-tuning steps, which were then used to generated test samples. We experimented on two trainable detectors, RoBERTa (semantic-based; 13) and SeqXGPT (probability-based; 43). Figure 2 displays **their decaying trends in accuracy as the fine-tuning corpus size increases**. For the LLM fine-tuned on 11.3M tokens, SeqXGPT, which performs almost *perfect* at detecting text pieces from the original LLM, even encounters an astonishing accuracy drop of 41%.

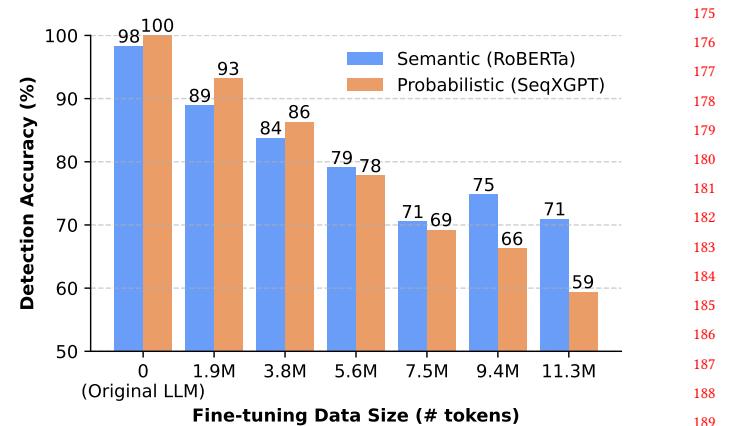


Figure 2: Detection accuracy for LLM-generated text with increasing amounts of fine-tuning data.

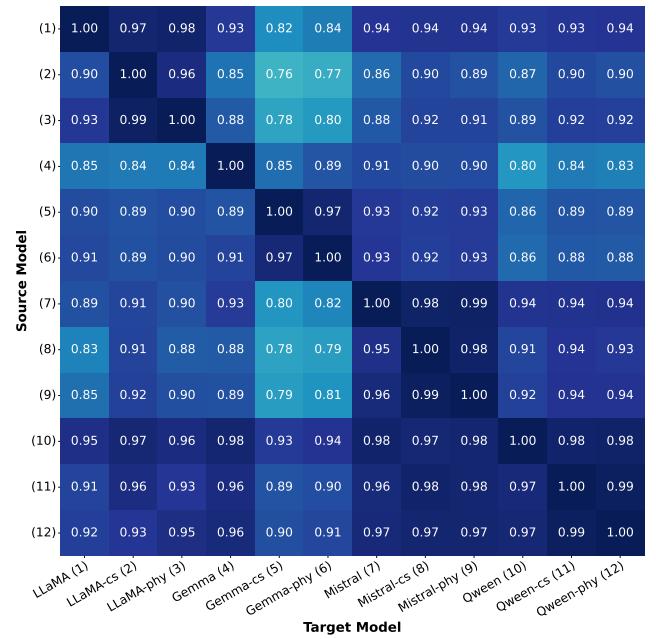


Figure 3: Normalized distance between the token probability lists of fine-tuned models and those of the base ones. Darker colors indicate higher similarity.

2.2 Impact of Fine-tuning on Features

Experimental Settings Our experiments are based on probability distributions from four open-source LLM families (LLaMA, Gemma, Mistral, and Qwen) across three domains (original, computer science, physics). We obtain probability distributions by processing JSONL files containing negative log-likelihood values and token ranks for each model-domain combination. Each file contains approximately 1,000 samples with complete probability distributions for analysis. Specifically, we collected responses from twelve model

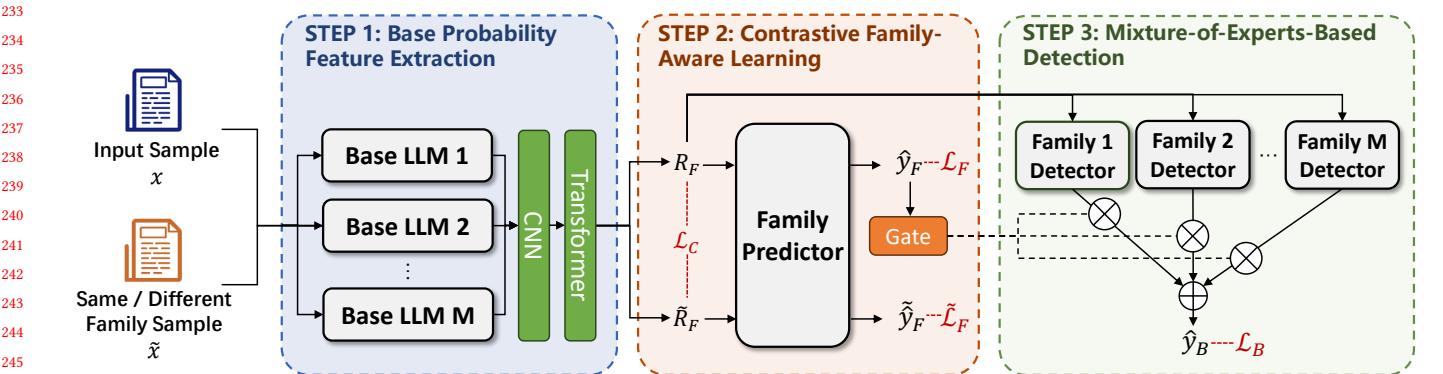


Figure 4: Overall architecture of PhantomHunter. Given a text sample x , it 1) extracts the probability feature from M base models and encode them with CNN and transformer blocks; 2) predicts the family of x to determine the family gating weights; and 3) feeds the representation R_F to a mixture-of-experts network controlled by the gating weights from Step 2 for final prediction of x being LLM-generated. During training, a contrastive learning is applied in each mini-batch to better model family relationships.

variants in total: four base models and twelve LoRA fine-tuned models (three domains \times four base models, see Dataset Construction section for detail). All responses were generated under identical decoding configurations (here, temperature = 1.0, top-p = 0.95) to ensure consistency.

Method. We implement a straightforward cross-model probability analysis to compute distribution similarities across model-domain combinations. For a given source model S and target model T , we extract their respective probability lists $P_S = [p_1, p_2, \dots, p_n]$ and $P_T = [q_1, q_2, \dots, q_n]$ from JSONL files containing negative log-likelihood values. To measure the divergence between probability distributions, we compute the normalized distance:

$$\text{dist}(P_S, P_T) = 1 - \frac{|\sum_{i=1}^n p_i - \sum_{i=1}^n q_i|}{\sum_{i=1}^n p_i}, \quad (1)$$

where the distance is bounded in $[0, 1]$ with higher values indicating greater similarity.

Findings and Inspiration. From Figure 3, we intuitively observe that the probability lists of LLMs from the same LLM are more similar to each other, and clearly different from those of the others, which shows significant “family trait”. Despite exceptions, an overall trend of family-wise LLM similarity is persisted. Exceptions may be caused by fine-tuning mechanics in different families and the coarse-grained metric. This observation reveals that the probability lists of the base models could be useful for modeling unseen fine-tuned ones. However, such “family traits” are not fully exploited by existing detectors because they treat all LLMs equally and may only learn the commonality across all involved LLMs (regarding the 12×12 area in the heatmap), not the family-level commonality (regarding each 3×3 areas), shaping a looser boundary against human text. Targetedly, we design PhantomHunter to learn such traits to address this issue.

3 Proposed Method: PhantomHunter

PhantomHunter is a family-aware learning-based detector for unseen privately-tuned LLM-generated texts. The key insight behind PhantomHunter is to model the inherent family-level features shared between base LLMs and their fine-tuned descendants. Building on these observations, PhantomHunter employs a family-aware learning approach to detect texts from unseen privately-tuned LLMs. As illustrated in Figure 4, the architecture consists of three main components: 1) a base probability feature extractor, 2) a contrastive learning-based family encoder, and 3) a mixture-of-experts detection module. We detail each component in the following subsections.

3.1 Base Probability Feature Extraction

For a given text sequence x , we pass it through M base LLMs $\theta_1, \theta_2, \dots, \theta_M$ to obtain the probability lists $p \in \mathbb{R}^{M \times N}$:

$$p = (p^{\theta_1}, p^{\theta_2}, \dots, p^{\theta_M}), \quad (2)$$

where $p^{\theta_i} = (p_1^{\theta_i}, p_2^{\theta_i}, \dots, p_N^{\theta_i})$, N is the text sequence length, M is the number of base models (*i.e.*, families), and $p_j^{\theta_i}$ represents the probability of generating token x_j given context $x_{<j}$ using LLM θ_i . We then use convolutional neural networks and Transformer encoders to extract base probability features $R_F \in \mathbb{R}^{M \times d}$, where d is the feature dimension.

3.2 Contrastive Family-Aware Learning

To better model family relationships in the feature space and enhance the detector’s generalization within the same family, we treat samples from the same family as augmentations and employ contrastive learning to enhance the representation R_F . Specifically, within each mini-batch, texts generated by models from the same family as the original text’s model are treated as positive examples, while others are negative ones. We compute a SimCLR-style [6]

349 contrastive loss \mathcal{L}_C :

$$351 \quad \mathcal{L}_C = -\sum_{i \in b} \log \frac{\exp(\delta(\mathbf{R}_{F_i}^m, \tilde{\mathbf{R}}_{F_i}^m)/t)}{\sum_{j=1}^{2|b|} \mathbf{1}_{[j \neq i]} \exp(\delta(\mathbf{R}_{F_i}^m, \mathbf{R}_{F_j}^m)/t)}, \quad (3)$$

354 where $\mathbf{R}_{F_i}^m/\tilde{\mathbf{R}}_{F_i}^m$ are the embeddings of the original/augmented sample,
355 respectively, δ is a dot-product function, and t controls the
356 temperature. We use a multi-layer perceptron (MLP) to classify the
357 LLM family:

$$359 \quad \hat{y}_F = \text{softmax}(\text{MLP}_F(\mathbf{R}_F)), \quad (4)$$

360 where $\hat{y}_F \in \{\theta_1, \theta_2, \dots, \theta_M\}$ represents the predicted LLM family.

362 3.3 Mixture-of-Experts-Based Detection

364 Finally, we employ a mixture-of-experts network for binary classification of human versus AI-generated text. Each expert detector
365 specializes in detecting texts from specific LLM families, enhancing
366 detection performance. Specifically, we use the family prediction
367 result \hat{y}_F as a gating signal to control the weighting of each expert's
368 prediction, yielding the final binary judgement:

$$371 \quad \hat{y}_B = \sum_{i=1}^M \hat{y}_F^i \cdot \text{softmax}(\text{MLP}_B(\mathbf{R}_F)), \quad (5)$$

374 where \hat{y}_F^i represents the probability that the text belongs to the
375 family of base model θ_i according to the family classifier. $\hat{y}_B \in [0, 1]$
376 denotes the probability of the x being LLM-generated. We optimize
377 the PhantomHunter model with:

$$379 \quad \mathcal{L} = \lambda_1(\mathcal{L}_F + \tilde{\mathcal{L}}_F) + \lambda_2 \mathcal{L}_B + \lambda_3 \mathcal{L}_C, \quad (6)$$

380 where \mathcal{L}_F and $\tilde{\mathcal{L}}_F$ represent the cross-entropy losses for family
381 classification of x and \tilde{x} respectively, \mathcal{L}_B is the cross-entropy loss
382 for LLM text detection, \mathcal{L}_C is the contrastive loss. λ_i ($i = 1, 2, 3$) are
383 hyperparameters to balance the loss items during the optimization.

385 PhantomHunter captures inherent signatures that persist through
386 fine-tuning by learning family-level features from observable fine-
387 tuned models derived from the same base LLMs, enabling detection
388 of unseen models from known families. Rather than focusing on in-
389 dividual model characteristics, we emphasize shared patterns defining
390 model families. This family-aware approach allows the detector
391 to generalize to unseen fine-tuned models by recognizing persis-
392 tent probability patterns that remain after extensive fine-tuning.
393 The method is effective because fine-tuning primarily adapts a
394 model's behavior to specific domains while preserving much of the
395 underlying probabilistic structure inherited from the base model.

396 4 Experiment

398 In this section, we conduct experiments to answer the following
399 questions:

400 **EQ1:** How effective is PhantomHunter at detecting texts from un-
401 seen fine-tuned LLMs?

402 **EQ2:** How do different components of PhantomHunter contribute
403 to its overall performance?

404 **EQ3:** How does PhantomHunter perform in diverse real-world
405 settings?

Table 1: Statistics of the corpora for fine-tuning.

ArXiv/Q&A Domain	# Tokens
Computer Science (cs)	6,441,516
Physics (phy)	5,955,389
Others (oth; q-bio/stat/eess/math/q-fin/econ)	3,939,626
ELI5	202,658
Finance (fin)	234,787
Medicine (med)	254,292

Table 2: Statistics of the evaluation dataset for privately-tuned LLMGT detection. FT Domain: Domains of data for LLM fine-tuning (“base” denotes no fine-tuning).

	Dataset	Source	FT Domain	#Samples
arXiv	Train	LLaMA	base/phy/oth	1,564/1,969/1,971
		Gemma	base/phy/oth	1,604/1,455/1,457
		Mistral	base/phy/oth	1,778/1,877/1,826
		Qwen	base/phy/oth	1,640/1,546/1,543
		Human	-	1,145
	Test	LLaMA	cs	2,076
		Gemma	cs	1,540
		Mistral	cs	2,002
		Qwen	cs	2560
		Human	-	208
Q&A	Train	LLaMA	base/med/ELI5	680/256/220
		Gemma	base/med/ELI5	680/256/220
		Mistral	base/med/ELI5	680/256/220
		Qwen	base/med/ELI5	680/256/220
		Human	-	5,780
	Test	LLaMA	fin	204
		Gemma	fin	204
		Mistral	fin	204
		Qwen	fin	204
		Human	-	806
	Total			32,818

4.1 Dataset Construction

We fully simulate the two most common usage scenarios of LLM: writing and question-answering. For writing, we collect 69,297 abstracts of academic papers from arXiv archive², categorizing them by primary subjects (domains). And for Q&A, we collect 3,062 Q&A pairs in ELI5, finance, and medicine domains from HC3 dataset [13]. Table 1 shows the corpora size. Note that we merge subjects other than physics and computer science for arXiv part to balance fine-tuning scales across domains. We select open-source LLaMA-2 7B-Chat [41], Gemma 7B-it [12], Mistral 7B-Instruct-v0.1 [21], and Qwen 2.5-7B-Instruct [2] as base models. For each scenario, we fine-tuned each base model using full-parameter fine-tuning (Full) and LoRA fine-tuning on the corresponding 3 domain-specific corpora from arxiv and Q&A, resulting in 48 derivative models. The details are as follows:

²<https://arxiv.org/archive/>

Table 3: F1 scores for human-written (Human), LLM-generated (Gen.), and both (Macro) and AUC for unseen fine-tuned LLM-generated text detection. The two best results in each metric are bolded and underlined, respectively. BFE: Base probability feature extraction. CL: Contrastive learning. MoE: Mixture-of-experts.

Method	arXiv								Q&A							
	Full				LoRA				Full				LoRA			
	Human	Gen.	Macro	AUC												
RoBERTa	62.42	98.54	80.48	97.63	65.61	98.66	82.13	93.16	84.30	95.12	89.71	98.72	58.12	78.03	68.07	99.75
T5-Sentinel	12.81	68.12	40.47	1.30	68.99	46.67	57.83	1.08	84.92	81.68	83.30	3.57	84.67	79.89	82.28	7.42
SeqXGPT	81.21	<u>99.52</u>	90.37	91.27	<u>93.15</u>	99.77	<u>96.46</u>	70.37	85.38	97.27	91.33	94.18	48.28	79.13	63.71	95.18
DNA-GPT	66.29	95.37	80.83	76.64	<u>72.08</u>	99.02	<u>85.55</u>	69.09	68.10	73.14	70.62	86.04	64.84	66.02	65.43	96.33
DetectGPT	26.71	95.50	61.11	85.56	52.53	98.33	<u>75.43</u>	79.27	59.46	88.72	74.09	82.80	26.09	<u>95.58</u>	61.11	92.36
Fast-DetectGPT	36.82	96.18	66.50	96.04	90.82	99.78	<u>95.30</u>	93.27	81.03	93.90	87.47	96.62	75.32	90.10	82.71	99.92
DALD	14.13	95.84	54.98	72.81	17.76	96.86	<u>57.31</u>	78.97	31.58	65.49	48.53	51.18	37.89	75.92	56.91	60.74
DeTeCtive	<u>85.09</u>	96.56	<u>90.83</u>	97.27	87.66	99.64	<u>93.65</u>	91.42	<u>89.08</u>	96.84	<u>92.96</u>	96.54	71.58	88.98	80.28	99.21
PhantomHunter	85.59	99.60	92.59	99.12	95.81	99.89	97.85	98.22	90.67	97.36	94.01	99.70	89.91	97.26	93.58	99.99
w/o BFE	66.45	98.74	82.59	99.58	67.54	98.77	83.16	99.73	81.60	<u>97.29</u>	89.45	99.77	69.86	87.91	78.89	99.84
w/o CL	78.29	99.31	88.80	96.34	92.44	<u>99.79</u>	96.12	96.84	81.97	94.33	88.15	99.89	82.82	95.08	88.95	99.99
w/o MoE	79.03	99.36	89.20	96.20	88.14	99.66	93.90	96.44	77.45	93.25	85.35	<u>99.80</u>	<u>83.05</u>	94.90	<u>88.97</u>	100.00

- **Full fine-tuning** was performed at the sft (supervised fine-tuning) stage. The training configuration was set with a batch size of 1 per device and gradient accumulation steps of 4. The learning rate was set to 1.0e-5, and the maximum number of training epochs was 10. A cosine learning rate scheduler was used with a warm-up ratio of 0.1. BF16 precision was enabled for efficient training, and the DDP (Distributed Data Parallel) strategy was employed.
- **LoRA fine-tuning** [19] was performed with the LoRA rank set to 8, and the training was conducted with a batch size of 1 per device, gradient accumulation steps of 4, a learning rate of 1.0e-5, and maximum number of epochs of 10, utilizing a cosine learning rate scheduler with a warm-up ratio of 0.1, BF16 precision for efficient training.

We then generate abstracts with 4 base models and 12 arXiv-based derivative models, and answers with the same base models and 12 Q&A-based derivative ones. For the arXiv dataset, we designate LLMs fine-tuned on computer science (cs) data as unseen ones, reserving them exclusively for testing. And for the Q&A dataset, we designate LLMs fine-tuned on finance (fin) data as unseen ones. Detailed statistics of our experimental dataset are shown in Table 2 and the adopted prompt templates are as follows:

Prompt Description: [Generating an abstract based on the title of an academic paper, used in the arxiv scenario]
Instruction Prompt: [Write an abstract for the academic paper titled *[title]*.]

Prompt Description: [Generating an answer based on the given question, used in the Q&A scenario]
Instruction Prompt: [*question*]

4.2 Experimental Setup

Implementation Details. In PhantomHunter, the CNN encoder comprises three convolutional layers followed by max-pooling, and the Transformer encoder has two attention layers with four attention heads each. The feature dimension d is set to 128. For contrastive learning, we use a temperature t of 0.07. The hyperparameters in the loss function are set as $\lambda_1 = 1.0$, $\lambda_2 = 1.0$, and $\lambda_3 = 0.5$, respectively. The model is trained using the Adam optimizer with a learning rate of 2e-5 and a batch size of 32. We train for 10 epochs and select the best checkpoint based on validation performance.

Evaluation Metrics. We report F1 scores on each testing subsets (i.e., human-written and LLM-generated subsets) with 0.5 as the threshold. We mainly analyze the results based on macro F1 scores as it provides a balance view. We also report the threshold-free metric AUC. For comparison with commercial detectors, we add the true positive rate under 1% false positive rate to better align with practical scenarios, following the recommendation by the recent literature Tufts et al. [42].

Compared Baselines. We select the following 8 methods for comparison: 1) **RoBERTa** [13]: A widely used detector based on a task-specific fine-tuned RoBERTa; 2) **T5-Sentinel** [7]: Another pretraining-based method for reframing the classification task as a next token prediction task; 3) **SeqXGPT** [43]: Also based on probability lists of the text, but reframes the classification task as a sequence labeling task; 4) **DNA-GPT** [48]: Determine the source of text based on multiple re-generations and statistical fingerprint analysis; 5) **DetectGPT** [31]: Determine whether a text is generated by comparing the probability of the original text piece with perturbed ones; 6) **Fast-DetectGPT** [3]: An efficient variant of DetectGPT that reduces the number of required perturbations; 7) **DALD** [51]: It aligns the surrogate model with the target model's distribution, thereby markedly enhancing curvature-and probability-divergence-based detection. 8) **DeTeCtive** [14]: A framework combining contextual cues and statistical features to

detect LLM-generated text. Due to the space limit, we detail the separate settings for all baselines in the supplementary material.

4.3 Main Results (EQ1)

Table 3 presents the binary classification results (human vs. LLM-generated) of PhantomHunter and baselines on unseen fine-tuned models. We have the following observations:

1) Compared to all baselines, PhantomHunter demonstrates superior performance in detecting texts from unseen fine-tuned models, with F1 scores for both human and LLM text exceeding 90% in each setting. For full fine-tuning, PhantomHunter improves the macro F1 score over the best baseline by 1.94% and 1.12% on both datasets, respectively; and for LoRA fine-tuning, the improvements are 1.44% and 5.18% respectively. This exhibits PhantomHunter’s detection capability for texts generated by unseen fine-tuned LLMs.

2) The fine-tuning approaches does not present consistent detectability for different domains or detectors. In most cases, the generated arXiv texts from LoRA-tuned LLMs are easier to detect than those from full-tuned ones; while the situations are reversed for Q&A texts. This might be because the arXiv texts are usually longer than Q&A texts, which could better reflect the effects of full parameter changes brought by full fine-tuning. In such comparison, PhantomHunter still keep balanced performance.

3) Some detectors show relatively low performance for the human class, regardless of class-wise balance of the training sets, indicating their difficulties in identifying human-written texts. We attribute this failure to their inherent ignorance of family traits shared among the base and derivative LLMs, which may shape a relatively blurred boundary between the human and LLM texts.

4.4 Ablation Studies (EQ2)

To evaluate the contribution of each component, we perform ablation studies by replacing base feature extractor to a RoBERTa (*w/o* BFE), removing contrastive loss (*w/o* CL), and removing mixture-of-experts components (*w/o* MoE) from PhantomHunter respectively, and the results are shown in Table 3.

From the results, we observe that removing any component leads to a degradation of detection performance, especially for BFE, since it is the probabilistic characterization of the base model that embodies the family commonality. Without contrastive loss, the probability representations of different family-generated texts become less distinguishable in the feature space. Due to the removal of the mixture-of-experts module, the classifier is forced to find a decision boundary between *all* family-generated texts and human texts without specializing in particular family models, which makes the learning process more difficult and thus lowers the detection performance.

4.5 Further Analysis (EQ3)

Analysis 1: Comparison with Commercial Detectors. We compare PhantomHunter with four commercial detectors that supports API calling on the whole test set, including WinstonAI³, BlueEyes⁴,

³<https://gowinston.ai/>

⁴[http://ruimeiruijianai.com/](http://ruimei.ruijianai.com/)

Table 4: Performance comparison with commercial detectors (%). *We remove samples with less than 300 characters for WinstonAI due to its minimum length constraint.

Method/Service	TPR _{1%FPR}	F1 _{Human}	F1 _{Gen}	F1 _{Macro}
WinstonAI*	32.81	27.48	87.29	57.39
Sapling	44.85	43.56	93.27	68.42
BlueEyes	59.26	35.85	89.24	62.55
HasteWire	69.90	65.01	97.50	81.26
PhantomHunter	99.45	93.43	99.56	96.50

Table 5: Multi-Domain Test Dataset Statistics from DetectRL

Domain	Source	Samples
ArXiv	Academic Abstracts	1,008
XSum	News Document Summaries	1,008
Writing Prompt	Creative Writing Stories	1,008
Yelp Review	Restaurant Reviews	1,008
Total		4,032

Sapling⁵, and HasteWire⁶. Due to the inaccessibility, we do not consider the commercial detectors with the requirement of additional actions besides API callings (e.g., Grammarly requires a subscription of the enterprise or education plan first) or have strict constraints in request times (e.g., Tencent’s Zhuque AI Detector allows only 20 requests per day). The details are as follows:

- **WinstonAI:** Due to its minimum length constraint of 300 characters, we remove shorter text samples for WinstonAI’s API requests, resulting in a smaller testing set of 6,492 samples (97.77% of the whole). Considering shorter text is generally harder to distinguish [40], the removal may bring WinstonAI detector an extra advantage over others.
- **BlueEyes:** A text auditing service launched by Ruijian AI, which supports LLM-generated text detection in Chinese and English. We access the API equipped with the latest version.
- **Sapling:** Sapling AI Content Detector uses a machine learning system (a Transformer network) similar to that used to generate AI content. The officially reported detection rate for AI-generated content is over 97% with less than 3% false positive rate for human-written content.
- **HasteWire:** An AI solutions provider for businesses that specializes in building custom AI solutions. It provides AI content detector which leverages advanced machine learning models to provide a confidence score on content origins. Its model has been trained against GPT-4o, Claude 3.5 Sonnet, DeepSeek-Chat, Llama-3.2, Mistral Small 2501, o3-mini and more.

We assume that commercial detectors have learned on text generated by the base models of involved families and been trained on a larger and more diverse training corpus than ours. However, Table 4 shows that PhantomHunter, trained on a small but focused dataset instead, significantly outperforms the compared ones, again highlighting the value of the proposed family-aware learning.

⁵<https://sapling.ai/>

⁶<https://hastewire.com/humanizer/detect>

Table 6: Generalization performance comparison among different methods (%), tested on the DetectRL benchmark. The two best results are bolded and underlined.

Method	macF1	F1Human	F1Gen
RoBERTa	39.36	00.00	78.71
T5-Sentinel	42.09	26.87	57.32
SeqXGPT	<u>74.15</u>	<u>81.04</u>	67.25
DetectGPT	58.41	61.90	54.91
DNA-GPT	64.18	67.58	60.77
FastDetectGPT	69.35	74.04	64.65
DeTeCtive	56.02	48.17	63.86
PhantomHunter	80.22	84.82	76.15
w/o BFE	63.70	56.21	71.20
w/o CL	71.37	75.45	67.29
w/o MoE	68.64	64.64	72.63

Analysis 2: Generalizability Test on the Public Benchmark

DetectRL. To test how the detectors generalizes to public benchmark data, we conduct *direct* inference of PhantomHunter and compared baselines trained on our corpus only against the representative public benchmark DetectRL [46], which covers diverse LLMs and multiple domains. As presented by Table 5, the setting encompasses 4 distinct domains, including academic abstracts from arXiv, news document summaries from XSum, creative writing stories from Writing Prompts, and restaurant reviews from Yelp. We followed the parameter configurations and feature extraction process from the main experiments described in our paper. We did not perform any training on this dataset and only conducted inference using the detectors trained on the data in the main experiment. From the results in Table 6, we see that PhantomHunter achieves a macro F1 score of 80.22% (84.82% for human-written text and 76.15% for LLM-generated text). Across the DetectEval multi-domain split (arXiv, WritingPrompts, XSum, and Yelp) it retains an average macro F1 score of 77.1%, confirming reliable cross-domain generalization without domain-specific tuning.

Analysis 3: A Practice of Including Seen LLMs. In practice, LLMGT detectors should maintain good performance against both seen and unseen LLMs. Here we show that PhantomHunter can be easily extend to be compatible to other seen LLMs through a simple modification. Specifically, an additional category “others” is added to the label space of the family classifier MLP_F , which is used as the family label of these texts from other seen models, and the number of experts in MoE is increased accordingly. Table 7 shows that PhantomHunter maintains high performance not only for unseen LLMs but also for seen GPT-4o mini⁷ and Claude 3.7 Sonnet⁸. Notably, we further introduced DeepSeek-V3⁹ and GLM-4-Air¹⁰ as entirely new LLM families unseen during training. Even under this setting, where CS data remained the unseen evaluation domain, the detectors continued to exhibit strong and stable performance, indicating its good compatibility and extendibility to more diverse LLMs.

⁷<https://openai.com/index/gpt-4o-mini-advancing-cost-efficient-intelligence/>

⁸<https://www.anthropic.com/news/clause-3-7-sonnet>

⁹<https://github.com/deepseek-ai/DeepSeek-V3>

¹⁰<https://github.com/zai-org/GLM-4>

Table 7: Performance comparison of seen and unseen fine-tuned LLM-generated texts (%), after adding the “others” family category.

	Model	TPR _{1%FPR}	F1Human	F1Gen	F1Macro
Unseen	LLaMA-cs	99.90	97.56	99.76	98.66
	Mistral-cs	99.90	98.04	99.80	98.92
	Gemma-cs	99.87	97.56	99.68	98.62
	Qwen-cs	98.98	93.02	99.41	96.22
	DeepSeek-v3	99.40	98.04	98.11	98.08
	GLM-4-air	99.77	98.04	98.11	98.07
Seen	GPT-4o mini	99.67	97.56	99.17	98.37
	Claude 3.7 Sonnet	100.00	98.04	99.50	98.77

Table 8: F1 scores for family classification of unseen fine-tuned LLM-generated texts.

	Dataset	LLaMA	Mistral	Gemma	Qwen	Macro
arXiv	Full	75.61	63.71	95.76	87.21	80.57
	LoRA	78.32	77.30	56.63	79.11	72.84
Q&A	Full	82.63	50.29	66.29	56.89	64.03
	LoRA	68.21	50.18	71.82	61.10	62.83

4.6 Exploration of Source Family Prediction

Results of Source Family Prediction. Beside enhancing the learning of LLMGT detection head, the introduction of a family classifier makes it possible to provide a source family prediction simultaneously, which facilitates further forensics analysis and may improve users’ trustworthiness to the detector. Table 8 previews the current performance of known family prediction. The performance is mixed. The highest F1 score exceeds 95% while the lowest is below 55%, demonstrating that the family prediction is far more challenging than LLMGT detection and the current design requires improvement for family prediction. Recalling the main results, this also implies that even a moderate family-aware learning could bring performance improvement, indicating that its potential remains to be further explored.

Case Analysis. To better understand the limitations of the current family classifier, we further examine misclassified samples using the *arXiv-Full* subset and analyze them from syntactic, lexical, stylistic, and structural perspectives. Our observations reveal several characteristic patterns:

1) The Gemma family exhibits highly distinctive and simplified syntactic behavior. Its outputs contain very few punctuation marks, shallow clause structures, and extremely low structural complexity (≈ 0.05), below LLaMA, Mistral, and human texts ($\approx 0.12\text{--}0.14$). Although Gemma shows unusually high lexical diversity ($TTR \approx 0.94$), it uses shorter words (average length ≈ 5.3). This extreme combination of “high diversity + short tokens” results in a unique feature island that the detector can recognize reliably.

2) Differently, the Qwen family unusually has long outputs (≈ 542 words) and prefers lexical reuse ($TTR \approx 0.327$), which serve as strong indicators. Even ignoring length, its moderate clause density and long-tail lexical statistics (e.g., modal density 0.22, verb-like

ratio 0.214) form a characteristic signature that is easy for the classifier to isolate.

3) For LLaMA and Mistral families, their syntactic density, structural complexity ($\approx 0.12\text{--}0.14$), and punctuation usage closely resemble human texts. This presents the most challenging cases. Their grammatical ratios—such as modal verbs, verb density, and pronoun usage—fall between each other and human writing, resulting in substantial overlap. Moreover, stylistic cues (e.g., word length, lexical richness) also align closely with human distributions. Because no single feature distinctly differentiates these models from humans, the classifier often misclassifies them.

Overall, Gemma and Qwen families form clearly separable “feature islands”, whereas LLaMA and Mistral cluster tightly with human writing across multiple linguistic dimensions, explaining why the classifier achieves higher accuracy on the former and often mispredicts the latter.

4.7 Cost Analysis

This section is to perform an analysis on the training and inference cost of PhantomHunter. We will show that despite the need of loading base LLMs, the overall computational cost of our method remains acceptable for real-world deployment.

- **On the inference side**, the system relies on four open-source models (Qwen2.5-7B, Gemma-7B-it, Mistral-7B-Instruct-v0.1, and LLaMA-2-7B-Chat) for probability feature extraction, each requiring roughly 14GB of GPU memory. The end-to-end latency is 37ms per input instance, including 28ms for sampling-based feature collection and 900ms for the classification step.
- **For training**, the model processes 1.15M tokens per epoch and completes an epoch in about 45s on an NVIDIA A800 GPU.

This demonstrates that our method achieves a significant improvement in accuracy at an acceptable cost. While the overall runtime remains practical, the reliance on multiple base models constitutes a cost bottleneck. In this work, we prioritize the possibility validation and would leaveganani this as a future research direction. The optimization may involve model distillation, feature sharing, or unified surrogate modeling.

5 Related Works

LLM-Generated Text Detection. Existing methods can be categorized as model-based and metric-based ones. The former generally trains neural models to analyze the textual characteristics of LLM outputs, exploiting semantic features [13, 14, 49] and self-familiarity [20]. The latter often sets similarity metrics based on LLMs’ output probability or behavioral characteristics [15, 31, 50, 52]. Recent methods collectively address different aspects of the LLM detection challenge, from data efficiency [53], to generalizability across models and domains [16], and training paradigms [45]. However, existing methods and evaluation benchmarks [17, 44, 46] focuses on publicly known LLMs. Such ignorance leads to detectors’ unsatisfying performance when countering texts from privately-tuned LLMs, as we experimentally showed before.

One possible solution is to add model-specific watermarks to the outputs as identifiers [27, 28], but this approach requires LLM providers to pay extra costs and recent research warns that LLM watermarks could be distorted or even removed through various

adversarial techniques [33, 47]. PhantomHunter showcases a non-watermarking solution to improve detection for text from unseen privately-tuned LLMs and it could work with watermarking solutions to more comprehensively lower the detection difficulty of text generated by privately-tuned LLMs.

Family Analysis of LLMs. Analyzing commonalities and uniqueness of LLMs from a family perspective yields a deep understanding of LLMs’ relationship [24], which is useful for LLM attributions [25, 38] and LLMs’ license violation detection [10]. As shown by Sarvazyan et al. [36] and Thorat and Yang [39], the family factor could influence text detectability and detectors’ generalizability. However, they only investigated the base models of different sizes, ignoring the unseen privately-tuned LLMs that we focus on in this research. In this work, we construct a new dataset containing texts generated by 48 LLMs that were privately fine-tuned with domain-specific corpus with two types of fine-tuning approaches. Our study is expected to provide new knowledge about the LLMGT’s detectability from the perspective of LLM post-training processes.

6 Conclusion and Discussions

We proposed PhantomHunter, which is specialized to improve detection performance of text generated by unseen privately-tuned LLMs. We first experimentally revealed that fine-tuning open-source base models would significantly lower detectability of generated text, and then addressed this issue via family-aware learning to capture shared traits in each LLM family. Comparison experiments with seven methods and four industrial services confirmed PhantomHunter’s superiority. The further analysis exhibited that PhantomHunter has good compatibility and extendibility to better adapt to real-world situations.

Discussions on Limitations. Though PhantomHunter showcases how to tackle the performance degradation issue when encountering texts generated by privately-tuned LLMs based on widely-known open-source ones, we also identify the following limitations of PhantomHunter which may influence the real-world implementation. First, the current version only supports binary classification between human-written and LLM-generated text plus the possible family source for a suspicious LLM-generated text piece. It does not provide fine-grained sentence- or token-level annotations [43] and the current performance of identifying family sources is moderate. Second, due to the experimental cost constraint, our experiment only covered the several common LLM families and the performance on other families like DeepSeek [8, 9] remains unknown. Third, extracting features from base LLMs requires them to be locally deployed and thus increases the memory cost and the inference computation overhead, which trades computation costs for higher detection performance.

Therefore, these directions requires further exploration and we advocate more industrial solutions to improve the trustworthiness and efficiency of LLMGT systems. Meanwhile, it is not recommended that the feedback of PhantomHunter serves as a sole basis for real-world actions (e.g., disciplinary action against any student). Additional harmful content detection [18, 26, 37] and human verification [32] are needed.

References

- [1] Alizila. 2025. Alibaba Unveils New Qwen3 Models for Coding, Complexing Reasoning and Machine Translation . <https://www.alizila.com/alibaba-unveils-new-qwen3-models-for-coding-complexing-reasoning-and-machine-translation/>. Accessed: 2025-08-02.
- [2] Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Bin Yuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chenggang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. 2023. Qwen Technical Report, arXiv:2309.16609 [cs.CL]
- [3] Guangsheng Bao, Yanbin Zhao, Zhiyang Teng, Linyi Yang, and Yue Zhang. 2024. Fast-DetectGPT: Efficient Zero-Shot Detection of Machine-Generated Text via Conditional Probability Curvature. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=BpcgrSE8Z>
- [4] Sid Black, Leo Gao, Phil Wang, Connor Leahy, and Stella Biderman. 2021. *GPT-Neo: Large Scale Autoregressive Language Modeling with Mesh-Tensorflow*. doi:10.5281/zenodo.5297715
- [5] Canyu Chen and Kai Shu. 2024. Combating Misinformation in the Age of LLMs: Opportunities and Challenges. *AI Magazine* 45, 3 (2024), 354–368. doi:10.1002/aaai.12188
- [6] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. A Simple Framework for Contrastive Learning of Visual Representations. In *Proceedings of the 37th International Conference on Machine Learning*. PMLR, 1597–1607. <https://proceedings.mlr.press/v119/chen20j.html>
- [7] Yutian Chen, Hao Kang, Vivian Zhai, Liangze Li, Rita Singh, and Bhiksha Raj. 2023. Token Prediction as Implicit Classification to Identify LLM-Generated Text. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 13112–13120. doi:10.18653/v1/2023.emnlp-main.810
- [8] DeepSeek-AI. 2025. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. arXiv:2501.12948 [cs.CL] <https://arxiv.org/abs/2501.12948>
- [9] DeepSeek-AI. 2025. DeepSeek-V3 Technical Report. arXiv:2412.19437 [cs.CL] <https://arxiv.org/abs/2412.19437>
- [10] Myles Foley, Ambrish Rawat, Taesung Lee, Yufang Hou, Gabriele Picco, and Giulio Zizzo. 2023. Matching Pairs: Attributing Fine-Tuned Models to their Pre-Trained Large Language Models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 7423–7442. doi:10.18653/v1/2023.acl-long.410
- [11] Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. SimCSE: Simple Contrastive Learning of Sentence Embeddings. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 6894. <https://aclanthology.org/2021.emnlp-main.552.pdf>
- [12] Gemma Team. 2024. Gemma: Open Models Based on Gemini Research and Technology. arXiv:2403.08295 [cs.CL] <https://arxiv.org/abs/2403.08295>
- [13] Biyang Guo, Xin Zhang, Ziyuan Wang, Minqi Jiang, Jinran Nie, Yuxuan Ding, Jianwei Yue, and Yupeng Wu. 2023. How Close is ChatGPT to Human Experts? Comparison Corpus, Evaluation, and Detection. arXiv:2301.07597 [cs.CL]
- [14] Xun Guo, Shan Zhang, Yongxin He, Ting Zhang, Wanquan Feng, Haibin Huang, and Chongyang Ma. 2024. DeTeCtive: Detecting AI-generated Text via Multi-Level Contrastive Learning. In *Advances in Neural Information Processing Systems*, Vol. 37. Curran Associates, Inc., 88320–88347. https://proceedings.neurips.cc/paper_files/paper/2024/file/a117a3cd54b7affad04618c77c2fb18b-Paper-Conference.pdf
- [15] Abhimanyu Hans, Avi Schwarzschild, Valeria Cherepanova, Hamid Kazemi, Aniruddha Saha, Micah Goldblum, Jonas Geiping, and Tom Goldstein. 2024. Spotting LLMs With Binoculars: Zero-Shot Detection of Machine-Generated Text. arXiv:2401.12010 [cs.CL]
- [16] Wei Hao, Ran Li, Weiliang Zhao, Junfeng Yang, and Chengzhi Mao. [n. d.]. Learning to Rewrite: Generalized LLM-Generated Text Detection. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 6421–6434. doi:10.18653/v1/2025.acl-long.322
- [17] Xinlei He, Xinyue Shen, Zeyuan Chen, Michael Backes, and Yang Zhang. 2024. MGTBench: Benchmarking Machine-Generated Text Detection. In *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*. ACM, 2251–2265. doi:10.1145/3658644.3670344
- [18] Beizhe Hu, Qiang Sheng, Juan Cao, Yuhui Shi, Yang Li, Danding Wang, and Peng Qi. 2024. Bad actor, good advisor: Exploring the role of large language models in fake news detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 22105–22113. doi:10.1609/aaai.v38i20.30214
- [19] Edward J Hu, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-Rank Adaptation of Large Language Models. In *The Tenth International Conference on Learning Representations*. <https://openreview.net/forum?id=nZeVKeeFYf9>
- [20] Jiazhou Ji, Jie Guo, Weidong Qiu, Zheng Huang, Yang Xu, Xinru Lu, Xiaoyu Jiang, Ruizhe Li, and Shujun Li. 2025. "I know myself better, but not really greatly": Using LLMs to Detect and Explain LLM-Generated Texts. arXiv:2502.12743 [cs.CL] <https://arxiv.org/abs/2502.12743>
- [21] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lamplé, Lucile Saulnier, Lélio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. 2023. Mistral 7B. arXiv:2310.06825 [cs.CL] <https://arxiv.org/abs/2310.06825>
- [22] Jean Kaddour, Joshua Harris, Maximilian Mozes, Herbie Bradley, Roberta Raileanu, and Robert McHardy. 2023. Challenges and Applications of Large Language Models. arXiv:2307.10169 [cs.CL] <https://arxiv.org/abs/2307.10169>
- [23] Ryuto Koike, Masahiro Kaneko, and Naoki Okazaki. 2024. OUTFOX: LLM-Generated Essay Detection Through In-Context Learning with Adversarially Generated Examples. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 21258–21266. doi:10.1609/aaai.v38i19.30120
- [24] Tharindu Kumarage, Garima Agrawal, Paras Sheth, Raha Moraffah, Aman Chadha, Joshua Garland, and Huan Liu. 2024. A Survey of AI-generated Text Forensic Systems: Detection, Attribution, and Characterization. arXiv:2403.01152 [cs.CL] <https://arxiv.org/abs/2403.01152>
- [25] Tharindu Kumarage and Huan Liu. 2023. Neural Authorship Attribution: Stylo-metric Analysis on Large Language Models. In *2023 International Conference on Cyber-Enabled Distributed Computing and Knowledge Discovery*. IEEE Computer Society, 51–54. doi:10.1109/CyberC58899.2023.00019
- [26] Yang Li, Qiang Sheng, Yehan Yang, Xueyao Zhang, and Juan Cao. 2025. From Judgment to Interference: Early Stopping LLM Harmful Outputs via Streaming Content Monitoring. *arXiv preprint arXiv:2506.09996* (2025).
- [27] Aiwei Liu, Leyi Pan, Yijian Lu, Jingjing Li, Xuming Hu, Xi Zhang, Lijie Wen, Irwin King, Hui Xiong, and Philip Yu. 2024. A Survey of Text Watermarking in the Era of Large Language Models. *Comput. Surveys* 57, 2 (2024), 1–36. doi:10.1145/3691626
- [28] Aiwei Liu, Qiang Sheng, and Xuming Hu. 2024. Preventing and Detecting Misinformation Generated by Large Language Models. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*. Association for Computing Machinery, 3001–3004. doi:10.1145/3626772.3661377
- [29] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. RoBERTa: A Robustly Optimized BERT Pretraining Approach. arXiv:1907.11692 [cs.CL] <https://arxiv.org/abs/1907.11692>
- [30] Meta AI. 2024. Llama Usage Doubled May Through July 2024. <https://ai.meta.com/blog/llama-usage-doubled-may-through-july-2024/>. Accessed: 2025-03-21.
- [31] Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D Manning, and Chelsea Finn. 2023. DetectGPT: Zero-Shot Machine-Generated Text Detection using Probability Curvature. In *Proceedings of the 40th International Conference on Machine Learning*. PMLR, 24950–24962. <https://proceedings.mlr.press/v2022/mitchell23a.html>
- [32] Chenhao Niu, Kevin P. Yancey, Ruidong Liu, Mirza Basim Baig, André Kenji Horie, and James Sharpnack. 2024. Detecting LLM-Assisted Cheating on Open-Ended Writing Tasks on Language Proficiency Tests. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track*. Association for Computational Linguistics, 940–953. doi:10.18653/v1/2024.emnlp-industry.70
- [33] Qi Pang, Shengyuan Hu, Wenting Zheng, and Virginia Smith. 2024. No Free Lunch in LLM Watermarking: Trade-offs in Watermarking Design Choices. In *Advances in Neural Information Processing Systems*, Vol. 37. Curran Associates, Inc., 138756–138788. https://proceedings.neurips.cc/paper_files/paper/2024/file/fa86a9c7b9f341716ccb679d1ae9fa-Paper-Conference.pdf
- [34] Giovanni Puccetti, Anna Rogers, Chiara Alzetta, Felice Dell'Orletta, and Andrea Esuli. 2024. AI 'News' Content Farms Are Easy to Make and Hard to Detect: A Case Study in Italian. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 15312–15338. doi:10.18653/v1/2024.acl-long.817
- [35] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. 2020. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *Journal of Machine Learning Research* 21 (2020), 1–67. <http://jmlr.org/papers/v21/20-074.html>
- [36] Areg Mikael Sarvazyan, José Ángel González, Paolo Rosso, and Marc Franco-Salvador. 2023. Supervised Machine-Generated Text Detectors: Family and Scale Matters. In *International Conference of the Cross-Language Evaluation Forum for European Languages*. Springer, 121–132. doi:10.1007/978-3-031-42448-9_11
- [37] Mrinank Sharma, Meg Tong, Jesse Mu, Jerry Wei, Jorrit Kruthoff, Scott Goodfriend, Euan Ong, Alwin Peng, Raj Agarwal, Cem Anil, Amanda Askell, Nathan Bailey, Joe Benton, Emma Bluemke, Samuel R. Bowman, Eric Christiansen, 9

987
988
989
990
991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1040
1041
1042
1043

- 1045 Hoagy Cunningham, Andy Dau, Anjali Gopal, Rob Gilson, Logan Graham, Logan Howard, Nimit Kalra, Taesung Lee, Kevin Lin, Peter Lofgren, Francesco
1046 Mosconi, Clare O'Hara, Catherine Olsson, Linda Petrini, Samir Rajani, Nikhil
1047 Saxena, Alex Silverstein, Tanya Singh, Theodore Sumers, Leonard Tang, Kevin K.
1048 Troy, Constantin Weisser, Ruiqi Zhong, Giulio Zhou, Jan Leike, Jared Kaplan,
1049 and Ethan Perez. 2025. Constitutional Classifiers: Defending against Universal
1050 Jailbreaks across Thousands of Hours of Red Teaming. arXiv:2501.18837 [cs.CL]
1051 <https://arxiv.org/abs/2501.18837>
- [38] Yuhui Shi, Qiang Sheng, Juan Cao, Hao Mi, Beizhe Hu, and Danding Wang. 2024. Ten Words Only Still Help: Improving Black-Box AI-Generated Text
1052 Detection via Proxy-Guided Efficient Re-Sampling. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence*. International Joint
1053 Conferences on Artificial Intelligence Organization, 494–502. doi:10.24963/ijcai.
1054 2024/55
- [39] Shantanu Thorat and Tianbao Yang. 2024. Which LLMs are Difficult to Detect?
1055 A Detailed Analysis of Potential Factors Contributing to Difficulties in LLM Text
1056 Detection. arXiv:2410.14875 [cs.CL] <https://arxiv.org/abs/2410.14875>
- [40] Yuchuan Tian, Hanting Chen, Xutao Wang, Zheyuan Bai, Qinghua Zhang,
1058 Ruifeng Li, Chao Xu, and Yunhe Wang. 2024. Multiscale Positive-Unlabeled
1059 Detection of AI-Generated Texts. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=5Lpq6U9hzV>
- [41] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillermo Cucurull, David Esioib, Jude Fernandes, Jeremy Fu, Wenying Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poultor, Jeremy Reizenstein, Rashi Rungra, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. Llama 2: Open Foundation and Fine-Tuned Chat Models. arXiv:2307.09288 [cs.CL]
1060 <https://arxiv.org/abs/2307.09288>
- [42] Brian Tufts, Xuandong Zhao, and Lei Li. 2025. A Practical Examination of AI-Generated Text Detectors for Large Language Models. arXiv:2412.05139 [cs.CL]
- [43] Pengyu Wang, Linyang Li, Ke Ren, Botian Jiang, Dong Zhang, and Xipeng Qiu. 2023. SeqXGPT: Sentence-Level AI-Generated Text Detection. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 1144–1156. <https://aclanthology.org/2023.emnlp-main.73/>
- [44] Yuxia Wang, Jonibek Mansurov, Petar Ivanov, Jinyan Su, Artem Shelmanov, Akim Tsvigun, Osama Mohanned Afzal, Tarek Mahmoud, Giovanni Puccetti, Thomas Arnold, et al. 2024. M4GT-Bench: Evaluation Benchmark for Black-Box Machine-Generated Text Detection. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 3964–3992. doi:10.18653/v1/2024.acl-long.218
- [45] Chenwang Wu, Yiu ming Cheung, Bo Han, and Defu Lian. 2025. Advancing Machine-Generated Text Detection from an Easy to Hard Supervision Perspective. arXiv:2511.00988 [cs.CL] <https://arxiv.org/abs/2511.00988>
- [46] Junchao Wu, Runzhe Zhan, Derek F. Wong, Shu Yang, Xinyi Yang, Yulin Yuan, and Lidia S. Chao. 2024. DetectRL: Benchmarking LLM-Generated Text Detection in Real-World Scenarios. In *Advances in Neural Information Processing Systems*, Vol. 37. Curran Associates, Inc., 100369–100401. https://proceedings.neurips.cc/paper/2024/file/b61bd7e9f64c04ec75a26e781e2ad51-Paper-Datasets_and_Benchmarks_Track.pdf
- [47] Qilong Wu and Varun Chandrasekaran. 2024. Bypassing LLM Watermarks with Color-Aware Substitutions. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 8549–8581. doi:10.18653/v1/2024.acl-long.464
- [48] Xianjun Yang, Wei Cheng, Yue Wu, Linda Ruth Petzold, William Yang Wang, and Haifeng Chen. 2024. DNA-GPT: Divergent N-Gram Analysis for Training-Free Detection of GPT-Generated Text. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=Xlayxj2fWp>
- [49] Xiao Yu, Kejiang Chen, Qi Yang, Weiming Zhang, and Nenghai Yu. 2024. Text Fluoroscopy: Detecting LLM-Generated Text through Intrinsic Features. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 15838–15846. doi:10.18653/v1/2024.emnlp-main.885
- [50] Xiao Yu, Yuang Qi, Kejiang Chen, Guoqiang Chen, Xi Yang, Pengyuan Zhu, Xiwei Shang, Weiming Zhang, and Nenghai Yu. 2024. DPIC: Decoupling Prompt and Intrinsic Characteristics for LLM Generated Text Detection. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*. <https://openreview.net/forum?id=BZh05P2EoN>
- [51] Cong Zeng, Shengkun Tang, Xianjun Yang, Yuanzhou Chen, Yiyou Sun, Zhiqiang Xu, Yao Li, Haifeng Chen, Wei Cheng, and Dongkuan DK Xu. 2024. DALD: Improving logits-based detector without logits from black-box llms. *Advances in Neural Information Processing Systems* 37 (2024), 54947–54973.
- [52] Biru Zhu, Lifan Yuan, Ganqu Cui, Yangyi Chen, Chong Fu, Bingxiang He, Yangdong Deng, Zhiyuan Liu, Maosong Sun, and Ming Gu. 2023. Beat LLMs at Their Own Game: Zero-Shot LLM-Generated Text Detection via Querying ChatGPT. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 7470–7483. doi:10.18653/v1/2023.emnlp-main.463
- [53] Xiaowei Zhu, Yubing Ren, Fang Fang, Qingfeng Tan, Shi Wang, and Yanan Cao. 2025. DNA-DetectLLM: Unveiling AI-Generated Text via a DNA-Inspired Mutation-Repair Paradigm. arXiv:2509.15550 [cs.CL] <https://arxiv.org/abs/2509.15550>

A Introduction of Compared Baselines

- **RoBERTa** [13]: A fine-tuned RoBERTa model [29] that uses semantic features to distinguish between human and AI-generated text.
- **T5-Sentinel** [7]: A model based on T5 [35] that treats token prediction as implicit classification to detect generated text.
- **SeqXGPT** [43]: A white-box method that treats token probability lists as waveforms and applies CNN and attention networks for detection. In the experiments, We use mistral-7b [21], Llama-2-7b-chat [41], Qwen2.5-7B-Instruct [2] and gemma-7b [12] as white-box models to compute the probability lists.
- **DNA-GPT** [48]: A source-attribution method that leverages multiple regenerations to identify model-specific statistical fingerprints, functioning in both black-box and white-box settings. In the experiment, we adopted the white-box setting and selected Llama-2-7b-chat [41] as the proxy base model, with regenerations num $N = 10$, the truncation ratio $\gamma = 0.5$, and max new tokens num $t = 250$. Since this method is training-free, we conducted a grid search for the classification threshold with a step size of $1e-4$, and reported the best result as the upper performance bound of this method.
- **DetectGPT** [31]: A zero-shot detection approach that identifies AI-generated text by comparing the probability of original passages against multiple perturbed variants without requiring model fine-tuning. In the experiment, we selected T5-3B [35] as the mask-filling model, and Llama-2-7b-chat [41] as the proxy base model, with the number of perturbations as 100. In experiment report, we follow a similar implementation to DNA-GPT.
- **Fast-DetectGPT** [3]: An efficient improvement over DetectGPT that introduces the concept of conditional probability curvature to identify machine-generated text, while significantly reducing computational overhead through optimized perturbation sampling strategies that maintain or enhance detection performance. In the experiment, we selected gpt-neo-2.7b [4] as the sampling model and the scoring model.
- **DALD** [51]: is a plug-and-play framework designed to enhance logits-based detectors in black-box settings. Unlike prior surrogate-based methods, DALD aligns the surrogate model's probability distribution with that of the unknown target model. The aligned surrogate model significantly improves curvature-based or probability-divergence-based detection metrics. Following the original paper's setup: Fine-tuning the surrogate model, Llama2-7B, on the WildChat dataset.

1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160

- 1161 • **DeTeCtive** [14]: A comprehensive framework that reformulates
 1162 the task from a binary classification problem to one of distin-
 1163 guishing various writing styles, treating each large language
 1164 model as a unique author with specific stylistic patterns. DeTeC-
 1165 tive combines multi-level contrastive learning and multi-task
 1166 auxiliary learning methods to enable fine-grained feature extrac-
 1167 tion that captures relationships between different text sources. In

the experiments, We use SimCSE-RoBERTa-base [11] as the pre-
 1219 train model. To preserve the original semantic space and prevent
 1220 overfitting, we freeze the model’s embedding layer during train-
 1221 ing. During testing, we adopt a KNN-based retrieval strategy
 1222 using the validation set to automatically determine the optimal
 1223 number of neighbors K within the range of 1 to 5, which is then
 1224 used for predicting out-of-distribution data.

1168

1225

1169

1226

1170

1227

1171

1228

1172

1229

1173

1230

1174

1231

1175

1232

1176

1233

1177

1234

1178

1235

1179

1236

1180

1237

1181

1238

1182

1239

1183

1240

1184

1241

1185

1242

1186

1243

1187

1244

1188

1245

1189

1246

1190

1247

1191

1248

1192

1249

1193

1250

1194

1251

1195

1252

1196

1253

1197

1254

1198

1255

1199

1256

1200

1257

1201

1258

1202

1259

1203

1260

1204

1261

1205

1262

1206

1263

1207

1264

1208

1265

1209

1266

1210

1267

1211

1268

1212

1269

1213

1270

1214

1271

1215

1272

1216

1273

1217

1274

1218

1275